

# CS 578

## Programming Assignment - 4.2

Mauli Mehulkumar Patel

Student ID: 35231636

### Design:

This is the design for the creation of a fault-tolerant system using zookeeper for the coordination and cassandra for the storage.

It attempts to implement checkpointing, fault tolerance, and recovery after failure.

### MyDBFaultTolerantServerZK.java :

#### Data outline:

- Client requests in /requests
- Server registers its address in /servers

#### Init:

- Create keyspace.
- Create grade and checkpoint tables.
- Load the last node from the checkpoint.

#### Architecture:

- Server: It is the key-value store. With zookeeper, it manages the clients. The leader will take in all the writes and connect to others, followers.
- Client: Sends the operations like GET etc.
- Zookeeper: It is the coordination service. It will keep a track of the current leader.

Leader Election used:

- Each server creates a node under servers which is created in the folder.
- Once numbers are given, the server with the lowest number becomes the leader.
- If a leader has failed, then the succeeding server becomes the leader after everyone is notified.

In case of Failure:

- Crash Recovery:
  - Server that recover reads from the stored /request
  - Compares the requests with its last znode
  - Covers up the missing operations

Checkpoint

- After operation, the last node is saved. Ensure constant overhead.

Concurrency:

- All servers accept clients parallel like threading.
- A leader change will call for some reconfiguration as only the leader can perform writes.

Failing Tests:

- Total Tests:
  - test31\_GracefulExecutionSingleRequest
  - test32\_GracefulExecutionMultipleRequestsSingleServer
  - test33\_GracefulExecutionMultipleRequestsToMultipleServers
  - test34\_SingleServerCrash
  - test35\_TwoServerCrash
  - test36\_OneServerRecoveryMultipleRequests
  - test37\_TwoServerRecoveryMultipleRequests
  - test38\_EntireStateMatchCheck
  - test39\_InstantaneousMassacreAndRevivalTest
  - test40\_SerialKillAndRecover
  - test41\_CheckpointRecoveryTest

- All the cases have passed

```

442, 1443, 1444, 1445, 1446, 1447, 1448, 1449, 1450, 1451, 1452, 1453, 1454, 1455, 1456, 1457, 1458, 1459, 1460, 1461, 1462, 1463, 1464, 1465, 1466, 1467, 1468, 1469, 1470, 1471, 1472, 1473, 1474, 1475, 1476, 1477, 1478, 1479, 1480, 1481, 1482, 1483, 1484, 1485, 1486, 1487, 1488, 1489, 1490, 1491, 1492, 1493, 1494, 1495, 1496, 1497, 1498, 1499, 1500, 1501, 1502, 1503, 1504, 1505, 1506, 1507, 1508, 1509, 1510, 1511, 1512, 1513, 1514, 1515, 1516, 1517, 1518, 1519, 1520, 1521, 1522, 1523, 1524, 1525, 1526, 1527, 1528, 1529, 1530, 1531, 1532, 1533, 1534, 1535, 1536, 1537, 1538, 1539, 1540, 1541, 1542, 1543, 1544, 1545, 1546, 1547, 1548, 1549, 1550, 1551, 1552, 1553, 1554, 1555, 1556, 1557, 1558, 1559, 1560, 1561, 1562, 1563, 1564, 1565, 1566, 1567, 1568, 1569, 1570, 1571, 1572, 1573, 1574, 1575, 1576, 1577, 1578, 1579, 1580, 1581, 1582, 1583, 1584, 1585, 1586, 1587, 1588, 1589, 1590, 1591, 1592, 1593, 1594, 1595, 1596, 1597, 1598, 1599, 1600, 1601, 1602, 1603, 1604, 1605, 1606, 1607, 1608, 1609, 1610, 1611, 1612, 1613, 1614, 1615, 1616, 1617, 1618, 1619, 1620, 1621, 1622, 1623, 1624, 1625, 1626, 1627, 1628, 1629, 1630, 1631, 1632, 1633, 1634, 1635, 1636, 1637, 1638, 1639, 1640, 1641, 1642, 1643, 1644, 1645, 1646, 1647, 1648, 1649, 1650, 1651, 1652, 1653, 1654, 1655, 1656, 1657, 1658, 1659, 1660, 1661, 1662, 1663, 1664, 1665, 1666, 1667, 1668, 1669, 1670, 1671, 1672, 1673, 1674, 1675, 1676, 1677, 1678, 1679, 1680, 1681, 1682, 1683, 1684, 1685, 1686, 1687, 1688, 1689, 1690, 1691, 1692, 1693, 1694, 1695, 1696, 1697, 1698, 1699, 1700, 1701, 1702, 1703, 1704, 1705, 1706, 1707, 1708, 1709, 1710, 1711, 1712, 1713, 1714, 1715, 1716, 1717, 1718, 1719, 1720, 1721, 1722, 1723, 1724, 1725, 1726, 1727, 1728, 1729, 1730, 1731, 1732, 1733, 1734, 1735, 1736, 1737, 1738, 1739, 1740, 1741, 1742, 1743, 1744, 1745, 1746, 1747, 1748, 1749, 1750, 1751, 1752, 1753, 1754, 1755, 1756, 1757, 1758, 1759, 1760, 1761, 1762, 1763, 1764, 1765, 1766, 1767, 1768, 1769, 1770, 1771, 1772, 1773, 1774, 1775, 1776, 1777, 1778, 1779, 1780, 1781, 1782, 1783, 1784, 1785, 1786, 1787, 1788, 1789, 1790, 1791, 1792, 1793, 1794, 1795, 1796, 1797, 1798, 1799, 1800, 1801, 1802, 1803, 1804, 1805, 1806, 1807, 1808, 1809, 1810, 1811, 1812, 1813, 1814, 1815, 1816, 1817, 1818, 1819, 1820, 1821, 1822, 1823, 1824, 1825, 1826, 1827, 1828, 1829, 1830, 1831, 1832, 1833, 1834, 1835, 1836, 1837, 1838, 1839, 1840, 1841, 1842, 1843, 1844, 1845, 1846, 1847, 1848, 1849, 1850, 1851, 1852, 1853, 1854, 1855, 1856, 1857, 1858, 1859, 1860, 1861, 1862, 1863, 1864, 1865, 1866, 1867, 1868, 1869, 1870, 1871, 1872, 1873, 1874, 1875, 1876, 1877, 1878, 1879, 1880, 1881, 1882, 1883, 1884, 1885, 1886, 1887, 1888, 1889, 1890, 1891, 1892, 1893, 1894, 1895, 1896, 1897, 1898, 1899, 1900, 1901, 1902, 1903, 1904, 1905, 1906, 1907, 1908, 1909, 1910, 1911, 1912, 1913, 1914, 1915, 1916, 1917, 1918, 1919, 1920, 1921, 1922, 1923, 1924, 1925, 1926, 1927, 1928, 1929, 1930, 1931, 1932, 1933, 1934, 1935, 1936, 1937, 1938, 1939, 1940, 1941, 1942, 1943, 1944, 1945, 1946, 1947, 1948, 1949, 1950, 1951, 1952, 1953, 1954, 1955, 1956, 1957, 1958, 1959, 1960, 1961, 1962, 1963, 1964, 1965, 1966, 1967, 1968, 1969, 1970, 1971, 1972, 1973, 1974, 1975, 1976, 1977, 1978, 1979, 1980, 1981, 1982, 1983, 1984, 1985, 1986, 1987, 1988, 1989, 1990, 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011, 2012, 2013, 2014, 2015, 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025, 2026, 2027, 2028, 2029, 2030, 2031, 2032, 2033, 2034, 2035, 2036, 2037, 2038, 2039, 2040, 2041, 2042, 2043, 2044, 2045, 2046, 2047, 2048, 2049, 2050, 2051, 2052, 2053, 2054, 2055, 2056, 2057], -1789828579=[29, 30, 32, 33, 36, 28, 34, 35, 31], -458500249=[0], -776943840=[13, 15, 14, 16, 18, 17], 1495874104=[31, 33, 28, 34, 29, 30, 32, 35, 36], 1468315478=[13, 15, 14, 16, 17, 18], -1961086961=[38, 41, 37, 42, 39, 40, 43, 44, 45], 774345459=[14, 13, 15, 17, 16, 18], 1028464473=[0], -1216554007=[8, 7, 9, 11, 12, 10], -174885816=[7, 8, 9, 10, 12, 11], 82292577=[19, 20, 21, 22, 23, 24, 25, 26, 27], -413574012=[1, 2, 3, 4, 5, 6], -895118981=[0], -331386271=[19, 20, 21, 22, 23, 24, 25, 26, 27], 1620811386=[9, 12, 8, 10, 11, 7], -1069325824=[7, 9, 11, 12, 8, 10], 1871757684=[7, 9, 10, 11, 12, 8], -1830260658=[19, 20, 21, 22, 23, 24, 25, 26, 27], 971440193=[0], -840343808=[37, 38, 41, 43, 44, 39, 40, 45, 42], -1499652239=[1, 2, 3, 4, 5, 6], -2140553852=[39, 45, 41, 43, 42, 44, 40, 37, 38], -1827320473=[19, 20, 21, 22, 23, 24, 25, 26, 27], 1387856779=[30, 33, 34, 36, 28, 29, 31, 35, 32], -159845200=[33, 31, 34, 35, 28, 30, 32, 36, 29], 1241122445=[39, 41, 40, 44, 42, 45, 43, 38, 37], -1951518250=[19, 20, 21, 22, 23, 24, 25, 26, 27], 2024844106=[28, 31, 33, 35, 36, 30, 34, 32, 29], -293600838=[1, 2, 3, 4, 5, 6], -1629604382=[1, 2, 3, 4, 5, 6], 2008210255=[0], 1425744005=[7, 10, 12, 9, 11, 8], -1925705529=[28, 30, 31, 34, 35, 29, 33, 36, 32]}
    succeeded
    .test49_DropTables succeeded
    .test99_closeSessionAndServers succeeded

Time: 162.442

OK (13 tests)

```

## Corrected failures:

- Failed Zookeeper locally: Initially there was some compatibility issue with the Mac and ZK
- Servers using wrong ports: Some servers were not connecting to correct ports and so clients were failing, and it indicated that connection refused.
- Restart test often failed due to empty grade table